

# OCR

January 1, 2023

```
[1]: !pip install --quiet -r requirements.txt
```

```
[2]: !nvidia-smi
```

Sun Jan 1 14:04:19 2023

```
+-----+
| NVIDIA-SMI 525.60.11      Driver Version: 525.60.11      CUDA Version: 12.0      |
+-----+-----+-----+-----+-----+-----+
| GPU  Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |                    MIG M. |
+=====+=====+=====+=====+=====+=====+
|   0  NVIDIA GeForce ...   Off   | 00000000:1F:00.0  On   |                     N/A |
| 83%   62C    P5    29W / 170W |   958MiB / 12288MiB |   16%      Default |
|                                       |                    |                    N/A |
+-----+-----+-----+-----+-----+-----+

```

```
+-----+
| Processes:
| GPU   GI    CI          PID    Type    Process name                        GPU Memory
|       ID    ID                                   Usage
+=====+
|    0   N/A   N/A         590      G   /usr/lib/Xorg                        403MiB |
|    0   N/A   N/A         627      G   /usr/bin/kwalletd5                   2MiB |
|    0   N/A   N/A         706      G   /usr/bin/ksmserver                    2MiB |
|    0   N/A   N/A         708      G   /usr/bin/kded5                        2MiB |
|    0   N/A   N/A         709      G   /usr/bin/kwin_x11                     83MiB |
|    0   N/A   N/A         755      G   /usr/bin/plasmashell                  48MiB |
|    0   N/A   N/A         772      G   ...de-authentication-agent-1          2MiB |
|    0   N/A   N/A         774      G   ...ib/xdg-desktop-portal-kde          30MiB |
|    0   N/A   N/A         854      G   /usr/bin/msm_kde_notifier              2MiB |
|    0   N/A   N/A         855      G   /usr/lib/kdeconnectd                  2MiB |
|    0   N/A   N/A         867      G   /usr/bin/kaccess                      2MiB |
|    0   N/A   N/A         870      G   /usr/bin/pamac-tray-plasma            2MiB |
|    0   N/A   N/A        1056      G   ...720691463375437107,131072         320MiB |
|    0   N/A   N/A        1261      G   ...-browser-integration-host         2MiB |
|    0   N/A   N/A        1270      G   ...AAAAAAAA= --shared-files           33MiB |
|    0   N/A   N/A        2101      G   /usr/bin/konsole                      2MiB |
+-----+

```

	0	N/A	N/A	3911	G	/usr/bin/dolphin	2MiB	
	0	N/A	N/A	6775	G	/usr/lib/kf5/kioslave5	2MiB	

-----

```
[3]: import os

filenames = []
for root, dirs, files in os.walk("CCPD2019-d11"):
    for f in files:
        if f.endswith(".jpg") or f.endswith(".JPG"):
            filenames.append(os.path.join(root, f))
filenames[:5]
```

```
[3]: ['CCPD2019-d11/test/0319- ASW872.jpg',
      'CCPD2019-d11/test/0401- AWW227.jpg',
      'CCPD2019-d11/test/0199- AXOL55.jpg',
      'CCPD2019-d11/test/0245- A9U395.jpg',
      'CCPD2019-d11/test/0389- AT219T.jpg']
```

```
[4]: len(filenames)
```

```
[4]: 209979
```

```
[5]: def get_label(filename: str) -> str:
      return filename.split(".")[0].split("-")[-1]

filenames[0], get_label(filenames[0])
```

```
[5]: ('CCPD2019-d11/test/0319- ASW872.jpg', ' ASW872')
```

```
[6]: import cv2
import matplotlib.pyplot as plt
from ipywidgets import interact, IntSlider

def plot(index: int):
    plt.imshow(cv2.cvtColor(cv2.imread(filenames[index]), cv2.COLOR_BGR2RGB));
    plt.axis("off")
    plt.show()

interact(plot, index=IntSlider(min=0, max=len(filenames) - 1, value=0));
```

```
interactive(children=(IntSlider(value=0, description='index', max=209978), Output()), _dom_classes=)
```

```
[7]: labels = list(map(get_label, filenames))
all_w, all_h, ratio = [], [], []

for filename in filenames[:20000]:
```

```
h, w, _ = cv2.imread(filename).shape
all_w.append(w)
all_h.append(h)
ratio.append(w / h)
```

```
[8]: import numpy as np

lens = list(map(len, labels))
np.mean(all_w), np.mean(all_h), np.mean(ratio), np.mean(lens)
```

```
[8]: (254.9241, 86.28205, 2.9521506126516366, 7.0)
```

```
[9]: from collections import Counter

print(Counter(all_w).most_common(5), Counter(all_h).most_common(5),
      ↪ Counter(ratio).most_common(5), Counter(lens).most_common(5), sep="\n")
```

```
[(244, 151), (225, 150), (235, 148), (231, 147), (220, 146)]
[(80, 504), (78, 494), (84, 492), (82, 492), (86, 491)]
[(3.0, 366), (2.75, 59), (2.8, 53), (2.6666666666666665, 53),
(2.857142857142857, 51)]
[(7, 209979)]
```

## 1 Augmentations

```
[10]: import albumentations as A

img_w, img_h = 256, 96

transforms = A.Compose([
    A.Resize(img_h, img_w, interpolation=cv2.INTER_CUBIC),
    A.Rotate(limit=10, p=0.25),
    A.RandomBrightnessContrast(brightness_limit=0.15, contrast_limit=0.15, p=0.
    ↪ 5)
])



val_transforms = A.Compose([
    A.Resize(img_h, img_w, interpolation=cv2.INTER_CUBIC),
])
```

```
[11]: img = cv2.cvtColor(cv2.imread(filenamees[0]), cv2.COLOR_BGR2RGB)
plt.imshow(img)
plt.axis("off");
```



```
[12]: nrows, ncols = 3, 3
_, axs = plt.subplots(nrows, ncols, figsize=(12, 12))

for i in range(nrows):
    for j in range(ncols):
        transformed = transforms(image=img)["image"]
        axs[i][j].imshow(transformed)
        axs[i][j].axis("off")
```



## 2 Tokenizer

```
[13]: import numpy as np

alphabet = np.array(sorted(list(set("".join(labels))))))
alphabet
```

```
[13]: array(['0', '1', '2', '3', '4', '5', '6', '7', '8', '9', 'A', 'B', 'C',
          'D', 'E', 'F', 'G', 'H', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q',
          'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z', ' ', ' ', ' ', ' ',
          ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ',
          ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ',
          ' '], dtype='<U1')
```

```
[14]: alphabet.shape
```

```
[14]: (66,)
```

```
[15]: !python --version
```

Python 3.10.8

```
[16]: #from __future__ import annotations

class Tokenizer:
    def __init__(self, alphabet: list[str]):
        self.word2token = {alphabet[i]: i for i in range(len(alphabet))}
        self.token2word = {v: k for k, v in self.word2token.items()}

    def encode(self, x: str) -> list[int]:
        return [self.word2token[i] for i in x]

    def __call__(self, x: str) -> list[int]:
        return self.encode(x)

    def decode(self, tokens: list[int]) -> str:
        return "".join([self.token2word[i] for i in tokens])

    def __len__(self) -> int:
        return len(self.word2token)

tokenizer = Tokenizer(alphabet)
len(tokenizer)
```

```
[16]: 66
```

```
[17]: tokenized = tokenizer(labels[0])
labels[0], tokenized, tokenizer.decode(tokenized)
```

```
[17]: (' ASW872', [51, 10, 27, 31, 8, 7, 2], ' ASW872')
```

### 3 Dataset

```
[18]: x_train = list(filter(lambda x: "train" in x, filenames))
x_test = list(filter(lambda x: "test" in x, filenames))
x_train[0], x_test[0]
```

```
[18]: ('CCPD2019-d11/train/0244025383142- APX882.jpg',
'CCPD2019-d11/test/0319- ASW872.jpg')
```

```
[19]: len(x_train), len(x_test)
```

```
[19]: (199980, 9999)
```

```
[20]: y_train = list(map(get_label, x_train))
      y_test = list(map(get_label, x_test))
      y_train[0], y_test[0]
```

```
[20]: (' APX882', ' ASW872')
```

```
[21]: len(y_train), len(y_test)
```

```
[21]: (199980, 9999)
```

```
[22]: import torch
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
torch.manual_seed(3407)
torch.cuda.manual_seed_all(3407)
device
```

```
[22]: device(type='cuda', index=0)
```

```
[23]: from torch.utils.data import Dataset, DataLoader
      from torchvision.transforms import ToTensor

      class OCRDataset(Dataset):
          def __init__(self, filenames: list[str], labels: list[str],
          →transforms=None):
              self.filenames, self.labels, self.transforms = filenames, labels,
          →transforms

          def __len__(self):
              return len(self.filenames)

          def __getitem__(self, index: int):
              img = cv2.cvtColor(cv2.imread(self.filenames[index]), cv2.COLOR_BGR2RGB)
              if self.transforms:
                  img = transforms(image=img)["image"]
              img = ToTensor()(img)
              return img, torch.LongTensor(tokenizer(self.labels[index]))

      train_ds = OCRDataset(x_train, y_train, transforms)
      val_ds = OCRDataset(x_test, y_test, val_transforms)

      batch_size = 128
      train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
          →num_workers=4)
      val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=False)

      train_ds[0]
```

```
[23]: (tensor([[[[0.0549, 0.0549, 0.0549, ..., 0.6196, 0.0549, 0.0745],
               [0.0549, 0.0549, 0.0549, ..., 0.6706, 0.0549, 0.0667],
               [0.0588, 0.0549, 0.0549, ..., 0.6706, 0.0549, 0.0627],
               ...,
               [0.1608, 0.1608, 0.1608, ..., 0.0549, 0.0549, 0.0549],
               [0.1765, 0.1804, 0.1765, ..., 0.0549, 0.0549, 0.0549],
               [0.1804, 0.1961, 0.2000, ..., 0.0549, 0.0549, 0.0549]],

               [[0.0627, 0.0588, 0.0588, ..., 0.6039, 0.0549, 0.0588],
               [0.0549, 0.0588, 0.0627, ..., 0.6667, 0.0549, 0.0549],
               [0.0549, 0.0549, 0.0549, ..., 0.6706, 0.0549, 0.0549],
               ...,
               [0.1686, 0.1647, 0.1608, ..., 0.0549, 0.0549, 0.0549],
               [0.1765, 0.1804, 0.1765, ..., 0.0549, 0.0549, 0.0549],
               [0.1804, 0.1961, 0.2000, ..., 0.0549, 0.0549, 0.0549]],

               [[0.0549, 0.0549, 0.0627, ..., 0.6000, 0.0549, 0.0549],
               [0.0549, 0.0549, 0.0588, ..., 0.6588, 0.0549, 0.0549],
               [0.0549, 0.0549, 0.0549, ..., 0.6706, 0.0549, 0.0588],
               ...,
               [0.1647, 0.1608, 0.1608, ..., 0.0549, 0.0549, 0.0549],
               [0.1686, 0.1725, 0.1765, ..., 0.0549, 0.0549, 0.0549],
               [0.1725, 0.1882, 0.2000, ..., 0.0549, 0.0549, 0.0549]]]],
        tensor([51, 10, 24, 32, 8, 8, 2])))
```

```
[24]: train_ds[0][0].size(), train_ds[0][1].size()
```

```
[24]: (torch.Size([3, 96, 256]), torch.Size([7]))
```

## 4 CRNN

```
[25]: import torch.nn as nn
import torch.nn.functional as F

class ConvBlock(nn.Module):
    def __init__(self, in_channels: int, out_channels: int, kernel_size=(3, 3),
        ↪ stride=1, padding=1):
        super().__init__()

        self.conv = nn.Conv2d(in_channels, out_channels, kernel_size, stride,
        ↪ padding)
        self.activate_fn = nn.ReLU(inplace=True)
        self.bn = nn.BatchNorm2d(out_channels)

    def forward(self, x):
```



```

        x = self.conv(x)
        x = self.activate_fn(x)
        x = self.bn(x)
        return x

class RNNBlock(nn.Module):
    def __init__(self, input_size: int, hidden_size: int, num_layers=2,
        ↳ dropout=0.1, bidirectional=True):
        super().__init__()

        self.rnn = nn.GRU(input_size, hidden_size, num_layers,
        ↳ batch_first=True, dropout=dropout, bidirectional=bidirectional)

    def forward(self, x):
        out, _ = self.rnn(x)
        return out

```

```

[26]: class CRNN(nn.Module):
    def __init__(self, num_classes: int):
        super().__init__()

        self.cnn = nn.Sequential(
            ConvBlock(3, 64),
            ConvBlock(64, 64),
            nn.MaxPool2d(kernel_size=2),

            ConvBlock(64, 128),
            ConvBlock(128, 128),
            nn.MaxPool2d(kernel_size=2),

            ConvBlock(128, 256),
            ConvBlock(256, 256),
            nn.MaxPool2d(kernel_size=2)
        )

        self.avg_pool = nn.AdaptiveAvgPool2d((7, 256))

        self.rnn = RNNBlock(256, 128)

        self.classifier = nn.Sequential(
            nn.Linear(128 * 2, 256),
            nn.ReLU(inplace=True),
            nn.Dropout(0.1),
            nn.Linear(256, num_classes)
        )

    def forward(self, x):

```

```

        x = self.cnn(x)
        bs, ch, h, w = x.size()
        x = x.reshape(bs, -1, 256)
        x = self.avg_pool(x)
        x = self.rnn(x)
        x = self.classifier(x)
        x = x.transpose(2, 1)
        return x

model = CRNN(len(tokenizer))
model = model.to(device)
model

```

```

[26]: CRNN(
  (cnn): Sequential(
    (0): ConvBlock(
      (conv): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (activate_fn): ReLU(inplace=True)
      (bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): ConvBlock(
      (conv): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (activate_fn): ReLU(inplace=True)
      (bn): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (3): ConvBlock(
      (conv): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (activate_fn): ReLU(inplace=True)
      (bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (4): ConvBlock(
      (conv): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      (activate_fn): ReLU(inplace=True)
      (bn): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (6): ConvBlock(
      (conv): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))

```

```

        (activate_fn): ReLU(inplace=True)
        (bn): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (7): ConvBlock(
        (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (activate_fn): ReLU(inplace=True)
        (bn): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (avg_pool): AdaptiveAvgPool2d(output_size=(7, 256))
    (rnn): RNNBlock(
        (rnn): GRU(256, 128, num_layers=2, batch_first=True, dropout=0.1,
bidirectional=True)
    )
    (classifier): Sequential(
        (0): Linear(in_features=256, out_features=256, bias=True)
        (1): ReLU(inplace=True)
        (2): Dropout(p=0.1, inplace=False)
        (3): Linear(in_features=256, out_features=66, bias=True)
    )
    )

```

```
[27]: from torchsummary import summary
```

```
summary(model, (3, 96, 256));
```

```

=====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
Sequential: 1-1                       [-1, 256, 12, 32]          --
|   ConvBlock: 2-1                     [-1, 64, 96, 256]          --
|   |   Conv2d: 3-1                     [-1, 64, 96, 256]          1,792
|   |   ReLU: 3-2                       [-1, 64, 96, 256]          --
|   |   BatchNorm2d: 3-3                 [-1, 64, 96, 256]          128
|   ConvBlock: 2-2                     [-1, 64, 96, 256]          --
|   |   Conv2d: 3-4                     [-1, 64, 96, 256]          36,928
|   |   ReLU: 3-5                       [-1, 64, 96, 256]          --
|   |   BatchNorm2d: 3-6                 [-1, 64, 96, 256]          128
|   MaxPool2d: 2-3                     [-1, 64, 48, 128]          --
|   ConvBlock: 2-4                     [-1, 128, 48, 128]         --

```

		Conv2d: 3-7	[-1, 128, 48, 128]	73,856
		ReLU: 3-8	[-1, 128, 48, 128]	--
		BatchNorm2d: 3-9	[-1, 128, 48, 128]	256
		ConvBlock: 2-5	[-1, 128, 48, 128]	--
		Conv2d: 3-10	[-1, 128, 48, 128]	147,584
		ReLU: 3-11	[-1, 128, 48, 128]	--
		BatchNorm2d: 3-12	[-1, 128, 48, 128]	256
		MaxPool2d: 2-6	[-1, 128, 24, 64]	--
		ConvBlock: 2-7	[-1, 256, 24, 64]	--
		Conv2d: 3-13	[-1, 256, 24, 64]	295,168
		ReLU: 3-14	[-1, 256, 24, 64]	--
		BatchNorm2d: 3-15	[-1, 256, 24, 64]	512
		ConvBlock: 2-8	[-1, 256, 24, 64]	--
		Conv2d: 3-16	[-1, 256, 24, 64]	590,080
		ReLU: 3-17	[-1, 256, 24, 64]	--
		BatchNorm2d: 3-18	[-1, 256, 24, 64]	512
		MaxPool2d: 2-9	[-1, 256, 12, 32]	--
		AdaptiveAvgPool2d: 1-2	[-1, 7, 256]	--
		RNNBlock: 1-3	[-1, 7, 256]	--
		GRU: 2-10	[-1, 7, 256]	592,896
		Sequential: 1-4	[-1, 7, 66]	--
		Linear: 2-11	[-1, 7, 256]	65,792
		ReLU: 2-12	[-1, 7, 256]	--
		Dropout: 2-13	[-1, 7, 256]	--
		Linear: 2-14	[-1, 7, 66]	16,962

=====

=====

Total params: 1,822,850  
Trainable params: 1,822,850  
Non-trainable params: 0  
Total mult-adds (G): 3.67

=====

=====

Input size (MB): 0.28  
Forward/backward pass size (MB): 84.03  
Params size (MB): 6.95  
Estimated Total Size (MB): 91.27

=====

=====

## 5 Utils

```
[28]: from Levenshtein import distance

def CER(pred: list[str], true: list[str]) -> float:
```

```

        return np.mean([distance(p, t) / max(len(p), len(t)) for p, t in zip(pred,
→true)])

def accuracy(pred: list[str], true: list[str]) -> float:
    return np.mean([p == t for p, t in zip(pred, true)])

CER([" APX882"], [" APX8782"]), CER([" APX882"], [" APX8772"]),
→accuracy([" APX882", " APX883", " APX884", " APX885"], [" APX882", " APX882",
→" APX884", " APX882"])

```

[28]: (0.125, 0.25, 0.5)

```

[29]: def prob2word(p: list[list[float]]) -> str:
        return tokenizer.decode([np.argmax(i) for i in p])

prob2word([[0.1, 0.6, 0.3, 0], [0.5, 0.4, 0.1, 0], [0, 0, 0, 1]],
→list(tokenizer.token2word.items())[:5]

```

[29]: ('103', [(0, '0'), (1, '1'), (2, '2'), (3, '3'), (4, '4')])

```

[30]: class AverageMeter(object):
        def __init__(self):
            self.reset()

        def reset(self):
            self.val = 0
            self.avg = 0
            self.sum = 0
            self.count = 0

        def update(self, val, n=1):
            self.val = val
            self.sum += val * n
            self.count += n
            self.avg = self.sum / self.count

```

## 6 Training

```

[31]: from tqdm.notebook import tqdm

def training(model, dl, criterion, optimizer, scheduler=None):
    model.train()
    loss_meter, acc_meter, cer_meter = AverageMeter(), AverageMeter(),
→AverageMeter()
    for x, y in tqdm(dl):

```

```

x, y = x.to(device), y.to(device)

optimizer.zero_grad()

out = model(x)

loss = criterion(out, y)

pred = list(map(prob2word, out.transpose(2, 1).cpu().detach().numpy()))
true = list(map(tokenizer.decode, y.cpu().detach().numpy()))

loss_meter.update(loss.item())
acc_meter.update(accuracy(pred, true))
cer_meter.update(CER(pred, true))

loss.backward()
optimizer.step()

if scheduler:
    scheduler.step()

del x
del y
del out
del loss
del pred
del true
torch.cuda.empty_cache()
return loss_meter.avg, acc_meter.avg, cer_meter.avg

def validation(model, dl, criterion):
    model.eval()
    loss_meter, acc_meter, cer_meter = AverageMeter(), AverageMeter(),
    ↪AverageMeter()
    for x, y in dl:
        x, y = x.to(device), y.to(device)

        with torch.no_grad():
            out = model(x)

        loss = criterion(out, y)

        pred = list(map(prob2word, out.transpose(2, 1).cpu().detach().numpy()))
        true = list(map(tokenizer.decode, y.cpu().detach().numpy()))

        loss_meter.update(loss.item())
        acc_meter.update(accuracy(pred, true))

```

```

        cer_meter.update(CER(pred, true))

    del x
    del y
    del out
    del loss
    del pred
    del true
    torch.cuda.empty_cache()
    return loss_meter.avg, acc_meter.avg, cer_meter.avg

```

```

[32]: import shutil
import gc

try:
    shutil.rmtree("checkpoints")
except FileNotFoundError:
    pass
finally:
    os.makedirs("checkpoints", exist_ok=True)

epochs = 20
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(model.parameters(), lr=3e-4)

best_val_loss = float("inf")

all_train_loss, all_train_acc, all_train_cer = [], [], []
all_val_loss, all_val_acc, all_val_cer = [], [], []

for e in range(1, epochs + 1):
    print(f"===== Epoch {e} / {epochs} =====")

    train_loss, train_acc, train_cer = training(model, train_dl, criterion,
    ↪optimizer)
    print(f"Train loss {train_loss} accuracy {train_acc} CER {train_cer}")
    all_train_loss.append(train_loss)
    all_train_acc.append(train_acc)
    all_train_cer.append(train_cer)

    val_loss, val_acc, val_cer = validation(model, val_dl, criterion)
    print(f"Validation loss {val_loss} accuracy {val_acc} CER {val_cer}\n")
    all_val_loss.append(val_loss)
    all_val_acc.append(val_acc)
    all_val_cer.append(val_cer)

    if val_loss < best_val_loss:

```

```
best_val_loss = val_loss
torch.save(model.state_dict(), f"checkpoints/epoch={e};loss={val_loss}")

gc.collect()
```

===== Epoch 1 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.45709893480896646 accuracy 0.6876967552201478 CER  
0.12333663617252884  
Validation loss 0.08500499303063637 accuracy 0.8981408227848101 CER  
0.02013166817359855

===== Epoch 2 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.026273303878983998 accuracy 0.9581293077124411 CER  
0.006739661490141335  
Validation loss 0.058788502706757076 accuracy 0.924940664556962 CER  
0.014508928571428567

===== Epoch 3 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.016522144474746373 accuracy 0.971769033909149 CER  
0.004424298510191025  
Validation loss 0.0422614835190905 accuracy 0.9452136075949367 CER  
0.01028481012658228

===== Epoch 4 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.011459875265659225 accuracy 0.9794670250974233 CER  
0.0031774247825111703  
Validation loss 0.04620456840560029 accuracy 0.9395239978902954 CER  
0.011267141350210974

===== Epoch 5 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.009397910206523705 accuracy 0.9829559086546851 CER  
0.0026633685013003764



Validation loss 0.045520578704102414 accuracy 0.9392800632911392 CER  
0.011556283905967453

===== Epoch 6 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.007390949275164066 accuracy 0.986434795410923 CER  
0.002113544444998388  
Validation loss 0.040412505729053215 accuracy 0.9461036392405063 CER  
0.01024242766726944

===== Epoch 7 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.005837200362419204 accuracy 0.9889489908683767 CER  
0.0016908225731402328  
Validation loss 0.03287486243356444 accuracy 0.9550039556962026 CER  
0.008236324593128393

===== Epoch 8 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.005147156110961811 accuracy 0.9905130358285349 CER  
0.0014752422082076556  
Validation loss 0.029045885817819757 accuracy 0.9597507911392406 CER  
0.00728978300180832

===== Epoch 9 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.004675342475732351 accuracy 0.9911778230966091 CER  
0.0013781304268348335  
Validation loss 0.029022535850932007 accuracy 0.9600474683544303 CER  
0.00721914556962025

===== Epoch 10 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.003920244600211253 accuracy 0.992522392834293 CER  
0.0011653413764738032  
Validation loss 0.02880717850164245 accuracy 0.9607397151898734 CER  
0.00720501808318264

===== Epoch 11 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.0036194277222167956 accuracy 0.9929627064793811 CER  
0.0010845879967760863  
Validation loss 0.023710776082110366 accuracy 0.9668710443037974 CER  
0.006060691681735983

===== Epoch 12 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.003245913009259259 accuracy 0.9936870201535508 CER  
0.0009725459281601239  
Validation loss 0.036751869217127184 accuracy 0.9495648734177216 CER  
0.009889240506329116

===== Epoch 13 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.002839836073081437 accuracy 0.9945721914441924 CER  
0.0008396663716961179  
Validation loss 0.025179882118289795 accuracy 0.9609375 CER 0.007021360759493665

===== Epoch 14 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.002806696706186022 accuracy 0.994476767434421 CER  
0.0008468718581482436  
Validation loss 0.026161488008359744 accuracy 0.9654865506329114 CER  
0.006244349005424954

===== Epoch 15 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

Train loss 0.0024493736921194784 accuracy 0.9953664827255279 CER  
0.0007090588154647622  
Validation loss 0.024989574671283814 accuracy 0.9648931962025317 CER  
0.006258476491862565

===== Epoch 16 / 20 =====

0%| | 0/1563 [00:00<?, ?it/s]

```
Train loss 0.00251026870657116 accuracy 0.9949625210841622 CER
0.000762483278078284
Validation loss 0.0282824625116714 accuracy 0.9642009493670886 CER
0.0065268987341772135
```

```
===== Epoch 17 / 20 =====
```

```
0%|          | 0/1563 [00:00<?, ?it/s]
```

```
Train loss 0.0021631120653590056 accuracy 0.995816338771593 CER
0.000637653093867105
Validation loss 0.02311562292448752 accuracy 0.9695411392405063 CER
0.005467337251356236
```

```
===== Epoch 18 / 20 =====
```

```
0%|          | 0/1563 [00:00<?, ?it/s]
```

```
Train loss 0.002014831066763753 accuracy 0.9961462332053743 CER
0.0005840988026688597
Validation loss 0.025750953161829633 accuracy 0.9669699367088608 CER
0.006131329113924047
```

```
===== Epoch 19 / 20 =====
```

```
0%|          | 0/1563 [00:00<?, ?it/s]
```

```
Train loss 0.00198260690326702 accuracy 0.9958713211772233 CER
0.000616231377387806
Validation loss 0.02216226708540351 accuracy 0.9716178797468354 CER
0.005142405063291135
```

```
===== Epoch 20 / 20 =====
```

```
0%|          | 0/1563 [00:00<?, ?it/s]
```

```
Train loss 0.0019587353731015285 accuracy 0.9961162428023033 CER
0.0005926674892605787
Validation loss 0.027547409252522963 accuracy 0.962618670886076 CER
0.006696428571428572
```


```
[33]: n = np.arange(1, epochs + 1)

plt.figure(figsize=(10, 10))
plt.plot(n, all_train_loss, label="Train loss")
```

```

plt.plot(n, all_train_acc, label="Train accuracy")
plt.plot(n, all_train_cer, label="Train CER")
plt.plot(n, all_val_loss, label="Validation loss")
plt.plot(n, all_val_acc, label="Validation accuracy")
plt.plot(n, all_val_cer, label="Validation CER")
plt.xlabel("Epoch")
plt.xticks(n)
plt.legend()
plt.show()

```



## 7 Evaluate

```
[34]: models = sorted(os.listdir("checkpoints"), reverse=True)
models
```

```
[34]: ['epoch=9;loss=0.029022535850932007',
'epoch=8;loss=0.029045885817819757',
'epoch=7;loss=0.03287486243356444',
'epoch=6;loss=0.040412505729053215',
'epoch=3;loss=0.0422614835190905',
'epoch=2;loss=0.058788502706757076',
'epoch=1;loss=0.08500499303063637',
'epoch=19;loss=0.02216226708540351',
'epoch=17;loss=0.02311562292448752',
'epoch=11;loss=0.023710776082110366',
'epoch=10;loss=0.02880717850164245']
```

```
[35]: from dataclasses import dataclass, field
from collections import defaultdict
from ipywidgets import Dropdown
import pandas as pd

@dataclass
class Info:
    filename: str = ""
    pred: str = ""
    true: str = ""
    metrics: list[float] = field(default_factory=list)

def evaluate(model_name):
    model = CRNN(len(tokenizer)).to(device)
    model.load_state_dict(torch.load(os.path.join("checkpoints", model_name)))
    model.eval()

    loss_meter, acc_meter, cer_meter = AverageMeter(), AverageMeter(),
↪AverageMeter()
    max_cer, max_loss = np.zeros(10), np.zeros(10)
    max_cer_info, max_loss_info = [Info() for _ in range(10)], [Info() for _ in
↪range(10)]
    for filename in tqdm(x_test):
        img = cv2.cvtColor(cv2.imread(filename), cv2.COLOR_BGR2RGB)
        x = val_transforms(image=img)["image"]
        x = torch.unsqueeze(torch.tensor(x), 0).to(device)
        label = get_label(filename)
        y = torch.unsqueeze(torch.tensor(tokenizer(label)), 0).to(device)

        with torch.no_grad():
```

```

        out = model(x)

        loss = criterion(out, y)

        pred = list(map(prob2word, out.transpose(2, 1).cpu().detach().
→numpy()))
        true = list(map(tokenizer.decode, y.cpu().detach().numpy()))

        cur_loss = loss.item()
        cur_acc = accuracy(pred, true)
        cur_cer = CER(pred, true)

        min_cer_index, min_loss_index = np.argmin(max_cer), np.
→argmin(max_loss)
        if cur_cer > max_cer[min_cer_index]:
            max_cer[min_cer_index] = cur_cer
            max_cer_info[min_cer_index] = Info(filename, pred[0], true[0],
→[cur_loss, cur_acc, cur_cer])
        if cur_loss > max_loss[min_loss_index]:
            max_loss[min_loss_index] = cur_loss
            max_loss_info[min_loss_index] = Info(filename, pred[0],
→true[0], [cur_loss, cur_acc, cur_cer])

        loss_meter.update(cur_loss)
        acc_meter.update(cur_acc)
        cer_meter.update(cur_cer)

        del x
        del y
        del out
        del loss
        del pred
        del true
        torch.cuda.empty_cache()

        max_cer, max_cer_info = zip(*sorted(zip(max_cer, max_cer_info), key=lambda
→x: -x[0]))
        max_loss, max_loss_info = zip(*sorted(zip(max_loss, max_loss_info),
→key=lambda x: -x[0]))

        cer_info = defaultdict(list)
        for d in list(map(lambda x: x.__dict__, max_cer_info)):
            for k, v in d.items():
                cer_info[k].append(v)

        loss_info = defaultdict(list)

```

```

for d in list(map(lambda x: x.__dict__, max_loss_info)):
    for k, v in d.items():
        loss_info[k].append(v)

pd.set_option("display.expand_frame_repr", False)

print(f"\n Test loss {loss_meter.avg} accuracy {acc_meter.avg} CER_{
↪cer_meter.avg}")
print("\n ==== Max CER files ==== \n")
print(pd.DataFrame(dict(cer_info), columns=["filename", "pred", "true", _
↪"metrics"]))
print("\n ==== Max Loss files ==== \n")
print(pd.DataFrame(dict(loss_info), columns=["filename", "pred", "true", _
↪"metrics"]))

@interact(index=IntSlider(min=0, max=len(x_test) - 1, value=0))
def check_model(index: int):
    img = cv2.cvtColor(cv2.imread(x_test[index]), cv2.COLOR_BGR2RGB)

    plt.imshow(img)
    plt.axis("off")
    plt.show()

    img = val_transforms(image=img)["image"]
    img = torch.unsqueeze(torch.tensor(img), 0).to(device)
    with torch.no_grad():
        out = model(img)
    out = out.squeeze()
    pred = prob2word(out.transpose(1, 0).cpu().detach().numpy())
    true = get_label(x_test[index])
    print("Predicted", pred)
    print("True", true)
    print("CER", CER([pred], [true]))

interact(evaluate, model_name=Dropdown(options=models));

```

```

interactive(children=(Dropdown(description='model_name', options=('epoch=9;loss=0.029022535850

```